

Build an Agentic Support Assistant

Yellow.ai · Delivery Team · Forward Deployed Engineer (Intern) · Screening Assignment

Time expected	1 day of focused work
Deadline	26th August 2026, Wednesday
Stack	Any language, framework, and LLM provider
Cost	Free LLM tiers only — do not spend money on this

The scenario

Trendly is a direct-to-consumer fashion retailer handling about 2,000 support chats a day. Roughly 70% are repetitive: order status, returns/exchanges, and shipping or refund policy questions. Your job as the FDE on this account is to build an agent that handles that volume end to end, and hands the rest to a human cleanly.

What we're assessing

- Orchestration — how the agent decides what to do, chains steps, recovers from failure, and carries state.
- Prompt engineering — how you control the model's behavior, ground it in the policy doc, and handle uncertainty.

What your agent should do

Across a multi-turn conversation, your agent should be able to:

- Look up an order and explain its status in plain language, including the edge cases.
- Answer policy questions grounded only in the provided policy document.
- Decide return/exchange eligibility by combining order data with policy rules, then act on it.
- Escalate to a human when it should, with a summary a person could actually use.
- Refuse what it shouldn't do — no invented policy, no unauthorized discounts, no data leakage.
- Also perform end-to-end testing and make sure it handles all edge cases and scenarios cleanly — no hallucinations.

Build this with real tool-calling (function calling, ReAct, a planner — your call), not keyword matching with an LLM wrapped around it. Design your own tools, guardrails, and approach to multi-turn state — we're deliberately not prescribing the shape, so use your own judgment.

What we provide

- `trendly_policy.md` — the shipping & returns policy. The only source of truth for policy questions.
- `orders.json` — 10 fixed orders. Load as-is.

Both files are available in this shared Drive folder:

<https://drive.google.com/drive/folders/1Okjfe445BstDkxid5gHS9W0XaPKtiUBf>

Deliverables

- Code — Git repo with a README (how to run it, base URL/start command, AI-usage note) and a PROMPTS.md documenting your prompts and how you iterated on them.
- Live endpoint — A live URL, or a repo that runs with one command, reachable for two weeks after the deadline.
- Solution note — 1–2 pages: architecture, key trade-offs, known limitations, and five discovery questions you'd ask Trendly's ops team before building this for real. Include this as a file in your repo (e.g. SOLUTION.md or a PDF) — the submission form does not have a separate field for it.
- Demo video — 3–5 minutes: happy path, two edge cases, one thing that doesn't work.

How to submit

When you're done, submit here:

<https://forms.gle/AmXBia95rouEXAks6>

You'll need: your Git repository URL, your live base URL (or start command), and a public demo video link.

How we evaluate

Scripted conversations testing order lookup & context, policy grounding, returns eligibility, escalation, safety & refusals, and robustness. Beyond that, there's no fixed spec — how you choose to design, extend, and reason through the ambiguity is part of what's being evaluated, not an extra.

Using AI tools

Use them freely — Claude, ChatGPT, Copilot, Cursor are all fine. Just note roughly what you generated vs. wrote in your README, and be ready to explain and modify your own code live if you're shortlisted.

Good luck — we want to see how you think.